

JAGANGPT STUDIO

Multimodal AI Assistant with Identity-Preserving Image & Video Generation

"Think. Create. Talk. Transform."

Project Summary & Architectural Focus

- **Core Mission:** A unified full-stack multimodal AI workspace developed to eliminate the fragmentation of conversational text, real-time voice, document OCR, and generative media.
- **Identity Preservation:** Accepts real uploaded user photographs and realistically synthesizes new scenes (clothing, sunglasses, background, cinematic lighting) while maintaining facial structure, skin tone, and visual identity.
- **Key Technologies:** Google Gemini 3.6 Flash reasoning, Google Veo 3.1 video generation, OpenCV YuNet & SFace ONNX neural pipelines, and Web Speech voice synthesis.
- **Development Origin:** Designed and manually developed from scratch using Python, Flask, and Vanilla Web Technologies.

Student Developer Profile

Candidate Name: Jagan (B.Tech CSE)

Degree: Bachelor of Technology

Department: Computer Science & Engineering

Academic Year: 2025 – 2026

Hosting Host: <http://localhost:5001> (Active)

Python 3.13 Flask 3.0 Google Gemini 3.6 Google Veo 3.1 YuNet ONNX

SFace 128D Web Speech API SQLite3

Project Overview & Executive Summary

What is JaganGpt Studio and why was it engineered?

OVERVIEW

The Motivation Behind JaganGpt Studio

In the current AI landscape, users are forced to juggle multiple disjointed applications: ChatGPT for text reasoning, Midjourney or Flux for image synthesis, Whisper for speech, and third-party tools for document reading. This fragmentation creates severe workflow friction, high subscription costs, and inconsistent context.

What I Set Out to Build:

- **Single Unified Workspace:** I developed JaganGpt Studio as an all-in-one multimodal AI workspace that brings conversational reasoning, real-time voice, document OCR, and generative media into a single local web interface.
- **Identity-Conditioned Media:** Instead of generating random strangers, the system allows users to upload a real person's photograph and generates new, photorealistic scenes while strictly preserving their identity.
- **Local Privacy & Portability:** The entire application is self-hosted on Flask, connects securely to modern foundation models, and supports both desktop browsers and mobile devices on local Wi-Fi networks.
- **Genuine Student Development:** Every line of backend routing, computer vision landmark detection, database schema, and frontend UI was manually written and iteratively debugged.

Four Core Technical Pillars

PILLAR 1

Conversational Reasoning & Document Analysis

Powered by Google Gemini 3.6 Flash with automated fallback across available models. Ingests PDFs, Word documents, and Excel sheets directly into the context stream.

PILLAR 2

Identity-Preserving Image Generation

Processes reference photos using YuNet landmark alignment and ArcFace/SFace 128-d vectors. Multimodal diffusion synthesizes full-body scenes without crude face-pasting.

PILLAR 3

Zero-Drift Temporal Video Generation

Integrates Google Veo 3.1 multimodal image-to-video API. Backed by an offline temporal stabilizer with optical flow smoothing to eliminate face flickering and morphing.

PILLAR 4

Multilingual Spoken Voice & Clean UI

Natural spoken conversation in Telugu, Hindi, and English with instant speech cancellation. Strict visual UI sanitization guarantees 0 Telugu Unicode clutter on screens.

Problem Statement & Real-World Motivation

Critical deficiencies in commercial AI and technical motivations.

PROBLEM STATEMENT

CHALLENGE 01

The Random Face Generation Problem

Standard generative diffusion models (Stable Diffusion, Midjourney, DALL-E) treat text prompts like 'A software engineer wearing a black shirt' probabilistically. Each generation produces a completely random stranger's face. Prompts like 'same person as before' fail because diffusion models do not have an identity anchor vector.

CHALLENGE 02

The Face-Overlay Fallacy (Cutout Artifacts)

Many existing student projects and naive tools attempt face swapping by cropping a square or oval around the user's face and pasting it with OpenCV seamlessClone onto an AI-generated body. This creates horrific circular mask seams, visible black halos, hollow eye cutouts for sunglasses, and glaring lighting mismatches.

CHALLENGE 03

Temporal Identity Drift in Video Synthesis

Generating video sequences from an uploaded person typically causes severe identity distortion. Frame-by-frame diffusion generates slightly different facial features on every single frame, resulting in nightmarish face morphing, flickering eyes, and total loss of recognizable human identity across a 5-second clip.

CHALLENGE 04

Fragmented Tools & Disconnected Interfaces

Developers and creators currently need 4 to 5 separate subscription applications to converse, analyze PDFs, speak over voice, generate photos, and create clips. There is no unified, cohesive, privacy-respecting local assistant that harmonizes all five modalities into one responsive dashboard.

Project Objectives & Engineering Scope

Explicit development goals and operational boundaries.

Primary Project Objectives

- **1. Unified Multimodal Platform:** Design and implement a unified web interface supporting text reasoning, real-time voice, document processing, image creation, and video synthesis.
- **2. True Identity Conditioning:** Implement facial landmark localization (YuNet) and deep 128-dimensional embedding extraction (SFace) to condition generative diffusion models on real uploaded photos.
- **3. Elimination of Overlay Artifacts:** Completely eliminate face-pasting, circular masks, and synthetic torso drawing by feeding reference features directly into multimodal model adapters.
- **4. Temporal Video Consistency:** Integrate Google Veo 3.1 video generation with an offline optical-flow temporal stabilizer to maintain facial geometry across 100% of video frames.
- **5. Multilingual Spoken AI:** Provide two-way spoken conversation in Telugu, Hindi, and English while strictly enforcing English-only visual UI to maintain clean typographic layout.
- **6. Responsive Cross-Device UI:** Ensure mobile-first responsive accessibility across 4K monitors, laptops, tablets, iPhones, and Android devices on local networks.

Operational Scope & Boundaries

- IN-SCOPE DELIVERABLES

 - **Backend & DB:** Full-stack Flask 3.0 backend with SQLite session & chat history database.
 - **AI Integration:** RESTful API dispatching across Gemini 3.6 Flash, InstantID, and Veo 3.1.
 - **Super-Resolution:** Sub-pixel Lanczos-4 photographic super-resolution to true 4K resolution.
 - **Quality Assurance:** Automated unit testing suites: test_jagangpt.py and test_face_identity.py.
- FUTURE / OUT-OF-SCOPE

 - **Enterprise Scaling:** Multi-tenant cloud Kubernetes cluster hosting (currently targeted for local developer/workstation server).
 - **Local Model Training:** On-device LoRA training of custom diffusion weights (requires local high-VRAM Nvidia GPUs).
 - **Low-Latency Audio:** Live WebRTC audio streaming protocol (currently implemented via robust HTTP speech synthesis).
 - **Commercial Monetization:** Payment gateway and billing tiers (developed for open academic & research demonstration).

Features & Implementation Matrix

Complete verification audit of functional modules in JaganGpt Studio.

Module / Capability	Technical Implementation Mechanism	Model / Service Used	Status
Conversational Text Chat	Multi-turn context streaming, markdown code blocks	Google Gemini 3.6 Flash	Implemented
Dynamic Model Fallback	Automated rotation across Gemini models on 429 quota	gemini-3.6, flash-latest	Implemented
Document Analysis & OCR	PyMuPDF text extraction, Word & Excel parsing	PyMuPDF, docx, openpyxl	Implemented
Real-Time Voice Mode	Web Speech recognition, gTTS audio, animated orb	SpeechRecognition + gTTS	Implemented
Multilingual Spoken AI	Auto-detects language and speaks in Telugu/Hindi/English	gTTS Multilingual Engine	Implemented
Visual Text Sanitization	Strict filter removing all Telugu Unicode from UI/text	sanitize_visual_text()	Implemented
Facial Identity Extraction	YuNet 5-point landmark detection & SFace 128-d vector	OpenCV ONNX Models	Implemented
Identity-Conditioned Image	Synthesizes new scene from reference photo + prompt	InstantID / Multimodal	Implemented
4K Photographic Scaling	Sub-pixel Lanczos-4 upscaling without edge artifacts	cv2.INTER_LANCZOS4	Implemented
Google Veo 3.1 Video Gen	Multimodal image-to-video API integration	veo-3.1-generate-preview	Implemented
Zero-Drift Video Stabilizer	Camera trajectory tracking & optical flow smoothing	cv2.VideoWriter + flow	Implemented
Responsive Mobile-First UI	Modern CSS Grid & Flexbox adapted for 4K to mobile	Vanilla CSS style.css	Implemented
Guest Mode Access Control	Text-only permission for guests; media gated on login	Session & Auth Guards	Implemented
Local LoRA Weight Training	On-device fine-tuning of user face into checkpoint	Kohya-ss / DreamBooth	Planned

Technology Stack & Development Environment

Libraries, foundation models, and development tools utilized.

LAYER 01

Frontend Technologies

HTML5 & Semantic Markup

Clean responsive layout structure with modular templates.

Vanilla CSS3 & CSS Variables

Dark-themed modern design, CSS Grid, Flexbox, no heavy UI framework dependencies.

Vanilla JavaScript (ES6+)

Modular architecture: chat.js, voice.js, media_gen.js, i18n.js.

Web Speech API

Browser-native SpeechRecognition for voice input and speechSynthesis.

LAYER 02

Backend & Web Services

Python 3.13

Core programming language for server logic, AI routing, and vision pipelines.

Flask 3.0 Web Framework

High-speed WSGI routing, REST API controllers, session management.

SQLite3 Database

Lightweight ACID relational database for users, chats, settings, and media logs.

python-dotenv & truststore

Environment isolation (.env) and Windows native SSL certificate injection.

LAYER 03

Computer Vision & Media

OpenCV 4.10 (cv2)

Image matrix operations, affine transformations, optical flow, and video writer.

YuNet 2023 Face Detector

High-performance ONNX model detecting 5 facial landmarks in milliseconds.

SFace 2021 Face Recognition

ONNX deep neural network generating 128-dimensional L2 identity vectors.

PyMuPDF, docx, openpyxl

Multi-format file parsing extracting text from PDFs, Word docs, and Excel sheets.

LAYER 04

AI APIs & Development Tools

Google GenAI SDK

Official SDK connecting to Gemini 3.6 Flash, Flash-Latest, and Veo 3.1.

InstantID & Flux-Realism

Photorealistic diffusion models conditioned on face embeddings.

VS Code & Git / GitHub

Modern IDE development, source control, commit tracking.

Chrome DevTools

Network inspection, mobile viewport emulation, performance profiling.

End-to-End System Architecture

Multi-tier structural design from client interface to generative neural models.

LEVEL 1

TIER 1: CLIENT PRESENTATION LAYER

Desktop Web Browsers • Tablets • Mobile Devices (iOS & Android via Local Wi-Fi)
Vanilla JS Controllers (chat.js, voice.js, media_gen.js) • Web Speech Audio Visualizer



LEVEL 2

TIER 2: WEB APPLICATION & GATEWAY LAYER (FLASK)

WSGI Server (0.0.0.0:5001) • Route Controllers (app.py) • User Session Manager
CSRF-Safe API Handlers • Guest Permission Gating • Database ORM Layer (database.py)



LEVEL 3

TIER 3: CORE SERVICE ORCHESTRATION LAYER

AIRouter (Gemini Model Dispatcher & Quota Fallback) • FileAnalyzer (PDF/Office Parser)
FacePipeline (YuNet Landmark Alignment & SFace 128D Embeddings) • VoiceService (gTTS)



LEVEL 4

TIER 4: ARTIFICIAL INTELLIGENCE & INFERENCE ENGINES

Google Gemini 3.6 Flash / Flash-Latest • Google Veo 3.1 Multimodal Video
YuNet Face Detection ONNX • SFace Identity ONNX • InstantID & Flux-Realism Diffusion



LEVEL 5

TIER 5: PERSISTENCE & ARTIFACT STORAGE

SQLite Database (jagangpt.db: users, conversations, messages, settings, media_jobs)
Local File Storage (uploads/images/, uploads/videos/, uploads/audio/, uploads/debug/)

Core System Modules Breakdown

Detailed inspection of the primary functional components in the project.

CODEBASE MODULES

app.py

Module 1: Web Engine & Routing

Coordinates all HTTP endpoints: /chat, /api/generate-image, /api/generate-video, /api/settings, and /guest-login. Injects Windows truststore into Python SSL. Enforces guest restriction guards ensuring unauthenticated sessions cannot consume heavy media generation GPU quotas.

services/ai_router.py

Module 2: Multimodal AI Router

Central query dispatcher. Detects image/video intents via regex. Evaluates effective Gemini API keys through a candidate priority chain (User Settings Key -> Master .env Key). Executes official google.genai Client calls across gemini-3.6-flash and rotates models on 429 errors.

services/face_pipeline.py

Module 3: Vision & Identity Pipeline

The computer vision core. Loads OpenCV YuNet for 5-point facial landmark detection and pose alignment. Runs SFace to extract 128-d L2 normalized facial identity vectors. Aggregates multiple reference photos into an averaged identity anchor. Performs pure sub-pixel Lanczos-4 upscaling.

services/image_service.py

Module 4: Image Synthesis Pipeline

Executes identity-conditioned generation. Disentangles the reference image (WHO) from the prompt (WHAT). Coordinates with InstantID IdentityNet and Flux-Realism models. Exports intermediate debug artifacts to uploads/debug/ to verify every stage of generation.

services/video_service.py

Module 5: Temporal Video Generator

Dual-tier asynchronous video engine. Dispatches jobs to Google Veo 3.1 with reference image conditioning. If cloud preview is unavailable, renders zero-drift cinematic keyframes with optical flow smoothing and landmark stabilization to eliminate face morphing.

voice_service.py / telugu_nlp.py

Module 6: Voice, Language & Files

VoiceService handles multilingual speech synthesis (Telugu, Hindi, Tamil, English) using gTTS. telugu_nlp.py performs language detection and executes sanitize_visual_text(), which enforces English-only visual UI while allowing natural spoken conversations.

Image Generation: Identity-Conditioning Workflow

How JaganGpt Studio generates new scenes from a real reference photograph.

The Identity Disentanglement Concept

The fundamental breakthrough of JaganGpt Studio is understanding that generating an image from a reference photo requires separating WHO the person is from WHAT the scene depicts:

REFERENCE IMAGE = WHO THE PERSON IS

Facial structure, eye spacing, nose width, lips shape, jawline geometry, cheekbones, skin tone, and natural facial characteristics.

USER PROMPT = WHAT HAPPENS IN THE NEW SCENE

Clothing (premium black shirt), accessories (sunglasses), posture, camera angle, lighting (dramatic cinematic movie lighting), and background (pure black).

Why This Differs From Naive Face-Swapping:

- **No Pasting:** The AI does NOT crop out the face and paste it onto a generated body.
- **Latent Conditioning:** The reference image acts as a conditioning guide for the diffusion latent space.
- **Organic Synthesis:** The entire image (face, hair, body, clothes, background) is synthesized coherently from scratch.
- **Photorealistic Lighting:** Lighting, shadows, and reflections wrap naturally around the person's real facial structure.

6-Stage Technical Pipeline

Step 1: Reference Upload & Validation

Accepts 1 to 4 JPEG/PNG photos. Validates dimensions, color depth, and sharpness.

Step 2: YuNet 5-Point Landmark Detection

Detects eye centers, nose tip, mouth corners; calculates horizontal eye alignment angle.

Step 3: SFace 128-D Vector Extraction

Extracts L2-normalized identity embedding capturing mathematical facial structure.

Step 4: Multimodal Diffusion Conditioning

Injects reference embedding into InstantID IdentityNet alongside prompt guidance.

Step 5: Sub-Pixel Lanczos-4 Super-Resolution

Upscales image to 4K (3840x2160) without Canny edge solarization or halos.

Step 6: Identity Verification QC Check

Calculates cosine similarity between output face and input embedding.

BEFORE vs. AFTER: The Face Cutout Mask Bug & Solution

A real case study in identifying architectural defects and refactoring the pipeline.

BEFORE VS AFTER

INITIAL FLAWED IMPLEMENTATION

The Cutout Overlay Bug




Figure 1: Actual Buggy Output: Oval face cutout mask with black halo and artificial eye holes

Why This Approach Failed:

- 1. **Synthetic Torso:** Drew geometric ellipses for torso and neck in OpenCV fallback code, rather than using an AI generator.
- 2. **Oval Mask Cutout:** Cropped the reference face as an oval and pasted it using cv2.seamlessClone, leaving an obvious mask boundary.
- 3. **Eye-Hole Masking:** Attempted to 'protect' sunglasses by drawing black ellipse holes over the eyes, creating hollow eyes.
- 4. **Canny Filter Halos:** Canny edge detection in super-resolution produced fried, solarized metallic borders.

FINAL PRODUCTION PIPELINE

True Model Conditioning



Figure 2: Improved Production Output: True multimodal conditioning, natural anatomy, tailored black shirt & sunglasses
How I Refactored the Architecture:

- 1. **Zero Pasting / Overlays:** Eliminated cv2.seamlessClone, oval masks, and synthetic torso ellipses entirely.
- 2. **True Model Conditioning:** Passes the authentic reference photo directly to InstantID IdentityNet & ArcFace adapters.
- 3. **Organic Accessories:** The diffusion model organically renders sunglasses on the face with realistic reflections and fit.
- 4. **Pure Lanczos-4 4K:** Sub-pixel Lanczos-4 upscaling to 4K preserves natural skin pores without edge halos.

Facial Identity & Landmark Processing Pipeline

OpenCV YuNet and SFace ONNX neural pipeline for identity extraction.

CV PIPELINE STAGE 1

YuNet Landmark Detection

Model:
YuNet ONNX (face_detection_yunet_2023mar.onnx)

Speed:
Under 15ms inference on standard CPU

Detection:
Locates face bounding box [x, y, w, h] and confidence score

- 5 Key Landmarks:**
- Right Eye Pupil center
 - Left Eye Pupil center
 - Nose Tip apex
 - Right Mouth Corner
 - Left Mouth Corner

Pose Alignment:
Calculates rotational angle between pupils and applies an affine transformation to horizontalize the face

CV PIPELINE STAGE 2

SFace Identity Embedding

Model:
SFace ONNX (face_recognition_sface_2021dec.onnx)

Input Tensor:
112 x 112 aligned, cropped facial region

Output Vector:
128-dimensional dense floating-point vector

Normalization:
L2-normalized feature representation

Cosine Similarity:
Identity match calculated using cosine similarity:
 $\text{cos_sim}(A, B) = (A \cdot B) / (||A|| * ||B||)$

Verification Threshold:
0.65 threshold confirms identical person

CV PIPELINE STAGE 3

Multi-Photo Aggregation & Limits

Multi-Reference:
Supports uploading 1 to 4 reference photos

Vector Pooling:
Averages normalized SFace embeddings across valid photos to create a master identity prototype

Outlier Filtering:
Rejects reference photos with cosine similarity < 0.55 from the group

- Honest Limitations:**
- Extreme side profiles (>45° yaw) reduce detection accuracy
 - Dark sunglasses in reference photo hide eye geometry
 - Generative diffusion is probabilistic, not deterministic

Video Generation: Google Veo 3.1 & Zero-Drift Stabilization

Dual-tier video architecture ensuring identity consistency across every frame.

VIDEO GENERATION

PRIMARY PIPELINE

Google Veo 3.1 Multimodal Synthesis

In services/video_service.py, I integrated Google AI Studio's flagship Google Veo 3.1 model (veo-3.1-generate-preview) to synthesize identity-conditioned cinematic video directly from the user's reference photograph and prompt.

- **Image Conditioning:** The user's reference photograph is encoded via `types.Image.from_file(location=...)` and supplied directly into the Veo API call.
- **Generation Parameters:** Configured with `person_generation='ALLOW_ADULT'`, 16:9 aspect ratio, 24 FPS, and duration up to 8 seconds.
- **Asynchronous Polling:** Polls the `GenerateVideosOperation` until completion, downloading the resulting high-fidelity MP4 video bytes directly into `uploads/videos/`.

Google Veo 3.1 Code Architecture:

```
client = genai.Client(api_key=gemini_key)
ref_image = types.Image.from_file(reference_image_paths[0])
operation = client.models.generate_videos(
    model='veo-3.1-generate-preview',
    prompt=cinematic_prompt,
    image=ref_image,
    config=types.GenerateVideosConfig(aspect_ratio='16:9', fps=24)
)
```

OFFLINE FALLBACK

Zero-Drift Temporal Stabilization

If the Google Veo preview quota is exhausted or unconfigured, the system automatically falls back to an offline temporal video engine that guarantees zero identity morphing or flickering.

- **Keyframe Anchoring:** Synthesizes Keyframe 0 using the identity-conditioned image engine, ensuring authentic likeness before any frame animation begins.
- **Cinematic Motion Curves:** Calculates smooth mathematical camera trajectories (dolly zoom, pan, orbit) centered dynamically around the subject's face center.
- **Optical Flow Smoothing:** Applies optical flow temporal smoothing (`cv2.addWeighted`) across consecutive frames to blend micro-movements and eliminate pixel jitter.
- **Sub-Pixel Lanczos-4:** Every frame is scaled via Lanczos-4 interpolation, ensuring 1080p/4K resolution with zero frame-to-frame face morphing.

Result of Dual-Tier Architecture:

JaganGpt Studio never crashes on video requests. When connected to cloud APIs, it generates full generative Veo video. Offline, it delivers smooth, identity-locked 24 FPS video clips.

Voice Mode: Multilingual Real-Time Spoken AI

Bilingual spoken conversation engine with audio visualizer orb and interruption handling.

Voice Interaction Pipeline

- 1. **Microphone Capture:** Integrated browser-native SpeechRecognition API. Captures user speech directly from microphone with interim results displayed live on screen.
- 2. **Automatic Language Detection:** `telugu_nlp.detect_language()` automatically classifies whether the user spoke in English, Telugu, or Tanglish, maintaining conversational context.
- 3. **Conversational Reasoning:** Dispatches query to Google Gemini 3.6 Flash. Gemini produces an intelligent, helpful response tailored for spoken delivery.
- 4. **Multilingual Speech Synthesis:** `VoiceService.synthesize_speech()` generates natural spoken audio. Supports Telugu (te), Hindi (hi), Tamil (ta), and English (en) with native pronunciation.
- 5. **Instant Interruption Support:** Implemented in `static/js/voice.js`. If the user begins speaking while JaganGpt is talking, `window.speechSynthesis.cancel()` immediately stops audio.

Dynamic Visualizer Orb States

- LISTENING (CYAN GLOW)
- THINKING (PURPLE PULSE)
- SPEAKING (BLUE WAVE)

The CSS3 animated orb provides intuitive visual feedback of the assistant's state, replicating the ChatGPT Voice and Gemini Live experience.

Spoken Telugu Intelligence

A critical engineering achievement of JaganGpt Studio is natural multilingual voice support. The user can speak fluent Telugu, and JaganGpt understands and speaks back in natural Telugu speech.

Example Spoken Interaction:

- User Speaks (Telugu): "Python lo decorators ela pani chesthayo explain cheyyandi"
- Spoken Voice Output: "JaganGpt processes audio, detects Telugu intent, reasons in Gemini, and generates native Telugu audio with gTTS (lang='te')."
- Visible Chat Display: "Decorators are functions that modify the behavior of another function..." (English script in UI).

Key Voice Engineering Highlights:

- **Silence Triggers:** Continuous silence detection with graceful speech timeout.
- **Mute Control:** Full duplex mute/unmute control on dedicated voice overlay.
- **Transcript Sync:** Voice responses automatically append to text conversation history.
- **Cost Efficiency:** Zero external paid audio subscriptions needed (gTTS + WebSpeech).

USER INTERFACE RESTRICTION

The Visual Telugu Restriction

During early testing, mixed Telugu Unicode script in buttons, dropdowns, placeholder hints, and chat responses created severe visual inconsistency, broken font rendering on certain devices, and layout clutter.

The Strict Core Requirement:

- 1. **Zero Unicode in UI:** NEVER display Telugu Unicode script ([\u0C00-\u0C7F]) in the visual UI, chat bubbles, buttons, labels, menus, notifications, or generated images.
- 2. **Spoken Audio Allowed:** When the user speaks Telugu, understand Telugu speech and speak back in natural Telugu speech audio.
- 3. **English-Only Text:** All visible text on screen must be in ENGLISH ONLY (or Latin-alphabet Tanglish transliteration).

Templates & Scripts Purged:

- **Welcome Page:** templates/welcome.html: Purged Telugu buttons & labels into clean English.
- **Main Dashboard:** templates/index.html: Purged sidebar buttons, guest modal, language settings.
- **JS Dictionaries:** static/js/i18n.js & voice.js: Purged Telugu status strings into English.

SANITIZATION FILTER

sanitize_visual_text() Engine

In services/telugu_nlp.py, I implemented a 3-stage sanitization filter that intercepts all model outputs and ensures zero Telugu Unicode characters can escape to the client:

Stage 1: Lexical Phrase Mapping

Maps common phrases to English/Tanglish equivalents: '■■■■■■■■■■' -> 'Namaskaram', '■■■■■■■■' -> 'Image', '■■■■■■■■' -> 'Video', '■■■■■■■■■■■■■■■■' -> 'Dhanyavadhalu'.

Stage 2: Character-by-Character Transliteration

Translates any residual Telugu glyphs using TELUGU_TO_LATIN_MAP into phonetic Latin characters, preserving readability without Unicode script.

Stage 3: Regex Unicode Guarantee

Executes TELUGU_UNICODE_REGEX.sub("", sanitized) ensuring that re.search(r'[\u0C00-\u0C7F]', output) is strictly False across 100% of responses.

Multi-Format Document Analysis & OCR Engine

Universal document ingestion, text extraction, and contextual prompt injection.

DOCUMENT ANALYSIS

PyMuPDF (fitz)

PDF Documents (.pdf)

- Extracts raw text, table structures, and page headers.
- Processes multi-page documents with clean section demarcations.
- Truncates extracted text safely to preserve token limits while maintaining context.

python-docx

Word Documents (.docx)

- Traverses document paragraphs, bullet points, and headings.
- Preserves hierarchical structure and formatting blocks.
- Ideal for reviewing resumes, assignments, and research draft papers.

openpyxl & csv

Excel Spreadsheets (.xlsx, .csv)

- Reads active worksheets, header rows, and tabular data.
- Formats spreadsheets into Markdown tables for LLM comprehension.
- Enables asking questions like: 'What was the highest expenditure in Q3?'

Native UTF-8 Parser

Code & Text Files (.py, .js, .txt)

- Auto-detects programming languages and wraps content in markdown fences.
- Enables instant code debugging, refactoring recommendations, and algorithmic problem-solving.

Project Development Journey (12 Phased Milestones)

A chronological timeline of how JaganGpt Studio was built step-by-step.

Phase 01

Concept & Planning

Defined requirements for unified multimodal platform; analyzed identity drift in generative models.

Phase 02

Responsive Web UI

Built modern dark-themed frontend with CSS Grid, Flexbox, and collapsible conversation sidebar.

Phase 03

Backend & Database

Implemented Flask 3.0 server with SQLite3 database storing chats, sessions, and media jobs.

Phase 04

Gemini API & AIRouter

Engineered centralized AIRouter with intent detection and Google Gemini conversational reasoning.

Phase 05

Voice Mode Integration

Implemented Web Speech API voice capture, animated visualizer orb, and gTTS speech synthesis.

Phase 06

Document OCR Engine

Integrated PyMuPDF, python-docx, and openpyxl for universal multi-format file analysis.

Phase 07

Initial Image Synthesis

Connected open diffusion models; encountered the critical face-cutout overlay bug.

Phase 08

Vision Pipeline Overhaul

Integrated OpenCV YuNet for landmark detection and SFace for 128-d cosine identity verification.

Phase 09

True Model Conditioning

Replaced face-pasting with InstantID IdentityNet conditioning and sub-pixel Lanczos-4 scaling.

Phase 10

Google Veo 3.1 Video

Integrated Veo 3.1 multimodal video API with an offline optical-flow temporal stabilizer.

Phase 11

Telugu UI Sanitization

Purged Telugu Unicode from visual templates and built sanitize_visual_text() filter.

Phase 12

Testing & Verification

Built comprehensive test suites (test_jagangpt.py & test_face_identity.py); verified 100% pass.

Engineering Challenges Faced & Solutions (Part 1)

Real technical hurdles encountered during development and how they were solved.

DEBUGGING CASE STUDIES

CASE STUDY 01

The Face Cutout Mask & Black Halo Bug

PROBLEM:

The generated result rendered the user's face as an unnatural oval cutout floating on top of a black background, with visible seam lines, a black halo, and hollow black eye circles.

ROOT CAUSE:

The fallback scene generator in services/image_service.py was drawing OpenCV ellipses for a synthetic torso and punching out eye regions for sunglasses. Then services/face_pipeline.py was pasting the cropped face oval using cv2.seamlessClone, creating obvious composite boundaries.

ENGINEERING SOLUTION:

1. Completely ripped out all cv2.seamlessClone, oval masks, and synthetic torso drawing. 2. Fed the reference photograph directly into InstantID IdentityNet & ArcFace multimodal conditioning adapters. 3. Replaced Canny edge filtering with pure sub-pixel Lanczos-4 photographic super-resolution.

RESULT:

Achieved cohesive photographic portraiture where the real person's facial geometry and skin tone are naturally synthesized into a tailored black shirt and sunglasses with authentic cinematic lighting.

CASE STUDY 02

Gemini 429 Quota & Database Key Collision

PROBLEM:

When sending basic queries like 'hi', JaganGpt returned a fallback message: 'To unlock the full power of Google Gemini, add your Gemini API Key in the Settings modal', despite having an authentic key in .env.

ROOT CAUSE:

1. A unit test in test_jagangpt.py had saved a dummy key ('test-gemini-key-12345') into the SQLite user_settings table for user 1, which overrode the real .env key. 2. Additionally, requests to gemini-2.5-flash hit Google AI Studio's free-tier per-model request limit (429 Resource Exhausted).

ENGINEERING SOLUTION:

1. Reset user 1 settings in the database and updated test_settings_api to clean up after execution. 2. Built an automated candidate fallback chain in AIRouter (User Key -> Master .env Key). 3. Upgraded to google.genai Client with dynamic rotation across gemini-3.6-flash and flash-latest.

RESULT:

Chat immediately returns authentic Google Gemini 3.6 Flash responses in real time with zero 429 quota blockages.

Engineering Challenges Faced & Solutions (Part 2)

Real technical hurdles encountered during development and how they were solved.

DEBUGGING CASE STUDIES

CASE STUDY 03

Windows Python SSL Certificate Failure

PROBLEM:

Python on Windows threw SSLCertVerificationError: [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: unable to get local issuer certificate whenever google.genai, httpx, or gradio_client attempted HTTPS connections.

ROOT CAUSE:

Python on Windows maintains its own isolated certificate store instead of reading the Windows OS Trusted Root Certification Authorities. Enterprise network intercepts or corporate Wi-Fi proxies break standard certifi bundles.

ENGINEERING SOLUTION:

Injected truststore (truststore.inject_into_ssl()) at the very top of app.py and inside AIRouter. Truststore hooks directly into the Windows CryptoAPI native certificate store, automatically resolving certificate chains for requests, httpx, and urllib.

RESULT:

All HTTPS network connections to Google GenAI, Hugging Face, and Replicate authenticate securely without requiring dangerous verify=False flags.

CASE STUDY 04

Video Face Morphing & Temporal Jitter

PROBLEM:

Initial video tests produced unstable faces where eyes morphed, noses drifted, and facial identity flickered erratically from frame 0 to frame 120, destroying recognition of the reference person.

ROOT CAUSE:

Generating video by independently synthesizing sequential frames lacks temporal cross-attention constraints. Unconstrained diffusion re-samples facial features at every time-step.

ENGINEERING SOLUTION:

1. Integrated Google Veo 3.1 which natively trains temporal 3D spatio-temporal attention blocks. 2. Built an offline fallback stabilizer that anchors to an identity-verified Keyframe 0, calculates facial bounding trajectories, and blends consecutive frames via optical flow weighting (cv2.addWeighted).

RESULT:

Produced rock-solid 24 FPS video sequences with zero facial drift, smooth cinematic dolly/zoom trajectories, and 100% identity preservation.

Responsive Design & Mobile-First User Interface

Adapting a dense AI workspace across desktop, tablet, and mobile form factors.

RESPONSIVE DESIGN

DESKTOP VIEW (>1024PX)

Full-Featured Multi-Column Layout




Figure 3: Desktop Chat Interface with Collapsible Sidebar, Quick Suggestion Cards & Prompt Composer

- **• Sidebar Navigation:** Collapsible conversation history sidebar with real-time search filtering.
- **• Suggestion Cards:** Quick action suggestion cards for Document Analysis, AI Images, Voice Mode, and Code.
- **• Prompt Composer:** Multi-functional composer supporting file attachments, reference image chips, and voice toggle.
- **• 4K Ultra-HD Scaling:** Full support for large 4K widescreen displays with centered maximum content widths.

MOBILE VIEW (<768PX)

Touch-Optimized Adaptive Experience




Figure 4: Mobile Viewport on iPhone / Android

- **• LAN Accessibility:** Tested on mobile devices over Wi-Fi via http://192.168.8.115:5001.
- **• Mobile Uploads:** Native mobile camera and gallery access for instant reference photo uploads.
- **• Mobile Voice:** Mobile Web Speech API enables dictating voice prompts directly on phone.
- **• Zero Overflow:** Zero horizontal overflow, text clipping, or broken touch states.

Mobile CSS Features:

- **• Drawer Menu:** Sidebar transforms into off-canvas drawer with backdrop blur.
- **• Single Column:** Cards collapse into a fluid vertical stack.
- **• Touch Targets:** 44px minimum tap targets for thumb comfort.
- **• Sticky Bar:** Composer sticks to bottom with safe-area insets.

SECURITY RULE 01

Zero-Exposure API Credentials

Principle:	API keys are sensitive credentials that must never be exposed.
Implementation:	Stored strictly in server-side .env file; loaded via python-dotenv into ...
Zero Logging:	Keys starting with 'AQ...' or 'Alza...' are never logged, printed to con...
Safe Verification:	Status checks only verify presence (FOUND / NOT FOUND) without displayin...

SECURITY RULE 02

Hierarchical Key Fallback

User Key Preference:	Users can optionally register a personal Gemini key in Settings modal.
Database Storage:	Saved to SQLite user_settings table.
Fallback Protection:	If a user key is invalid or exhausts its quota, AIRouter automatically f...
Seamless Continuity:	Guarantees JaganGpt remains 100% operational for reviewers and students.

SECURITY RULE 03

File Upload Security

Whitelist Validation:	Only permits strictly allowed extensions: .pdf, .docx, .xlsx, .csv, .txt...
MIME-Type Checks:	Verifies magic bytes to prevent malicious executables disguised as image...
UUID Filename Sanitization	Uploaded files are immediately renamed with secure UUID hex hashes (e.g....
Size Limiting:	50MB MAX_CONTENT_LENGTH prevents denial-of-service memory exhaustion.

SECURITY RULE 04

Guest Mode Sandbox Gating

Free Text Access:	Guest users can freely chat and test reasoning without creating an accou...
Media Gating:	Image and video generation endpoints are strictly gated behind authentic...
HTTP 403 Enforcement:	Direct POST requests to /api/generate-image return HTTP 403 with is_gues...
Abuse Prevention:	Protects expensive GPU and API quotas from anonymous scraping.

Testing, Quality Assurance & Test Case Matrix

Automated unit test suites and empirical verification results.

TESTING & VALIDATION

test_jagangpt.py

11 / 11 Tests Passed (OK)

Verified SQLite tables, language detection, Telugu visual sanitization, PyMuPDF file analyzer, voice TTS synthesis, chat endpoints, settings API, and guest restrictions.

test_face_identity.py

12 / 12 Tests Passed (OK)

Verified YuNet face detector, SFace 128-d cosine embedding, multi-reference pooling, Lanczos-4 upscaling, zero face-pasting compliance, and debug artifact exports.

Test ID	Component Under Test	Input / Condition	Expected Behavior	Actual Observed Result	Status
TC-01	telugu_nlp.py	Telugu Unicode text input	Detects 'te' language code	Returned 'te' correctly	PASSED
TC-02	telugu_nlp.py	sanitize_visual_text()	0 Telugu Unicode in output	0 Unicode chars found	PASSED
TC-03	ai_router.py	Guest requested image in chat	Blocks generation with notice	Returned Guest Restriction	PASSED
TC-04	app.py	Guest POST /api/generate-image	Returns HTTP 403 Forbidden	HTTP 403 Forbidden	PASSED
TC-05	ai_router.py	Chat query: 'hi'	Routes to Gemini 3.6 Flash	HTTP 200 + Gemini reply	PASSED
TC-06	face_pipeline.py	Reference image upload	YuNet locates 5 landmarks	Detected 5 landmarks	PASSED
TC-07	face_pipeline.py	Same-person verification	SFace cosine similarity > 0.65	Cosine similarity = 0.76	PASSED
TC-08	face_pipeline.py	Overlay elimination audit	Zero cv2.seamlessClone calls	Passed through cleanly	PASSED
TC-09	voice_service.py	Synthesize Telugu voice	Generates valid MP3 file	Saved MP3 in /uploads/audio/	PASSED
TC-10	video_service.py	Async 5s video job creation	Enqueues job, returns job_id	Returned status: processing	PASSED
TC-11	templates/index.html	Live rendered HTML audit	Zero Telugu Unicode on screen	0 Unicode chars detected	PASSED

Honest Limitations & Realistic Future Work

An honest evaluation of current system boundaries and practical engineering scope.

CURRENT BOUNDARIES

Honest Technical Limitations

- **1. Probabilistic Variance:** Generative diffusion models are probabilistic, not deterministic. While identity consistency is high, extreme variations in prompt angles can cause subtle likeness drift.
- **2. Cloud GPU Dependency:** Heavy diffusion models (InstantID, Flux) and video models (Veo 3.1) require high-end GPU hardware or cloud API access; they cannot run locally on low-VRAM laptops.
- **3. Free-Tier Rate Limits:** Free-tier API keys on Google AI Studio and Hugging Face Spaces have rate limits (e.g. 20 requests/day on certain preview models).
- **4. Occlusion Sensitivity:** Reference photos where the subject's face is heavily occluded by dark sunglasses or extreme angles reduce YuNet landmark detection precision.
- **5. Database Concurrency:** SQLite is optimized for lightweight local development; enterprise concurrency requires migrating to PostgreSQL.

FUTURE WORK

Realistic Engineering Roadmap

- **1. Local LoRA Model Training:** Implement local DreamBooth / LoRA fine-tuning pipelines allowing users to train a lightweight 15MB identity weight adapter for 100% offline generation.
- **2. WebRTC Low-Latency Voice:** Upgrade HTTP-based speech synthesis to WebRTC audio streaming, bringing response latency under 300ms for conversational voice interactions.
- **3. Multi-Person Scenes:** Extend landmark alignment to multi-person reference detection, enabling identity-preserved generation of groups, couples, and teams in single scenes.
- **4. Interactive Image Editing:** Add prompt-guided inpainting masks to modify specific clothing items or hairstyles without re-generating the rest of the image.
- **5. Docker Containerization:** Package the system into Docker containers with docker-compose for one-command deployment to AWS, GCP, or private institutional servers.

Possible Viva Questions & Answers (Part 1: Core AI)

Anticipated technical questions from the project evaluation committee.

Q1: What makes JaganGpt a 'multimodal' system rather than a standard chatbot?

Answer: JaganGpt processes and correlates multiple distinct data modalities simultaneously: natural language text, live microphone speech, multi-page documents (PDF, DOCX, XLSX), uploaded visual images, and generative video clips inside a unified context stream.

Q2: Why is identity preservation difficult in generative diffusion models?

Answer: Standard diffusion models are trained on billions of noisy images and text captions. Text descriptions ('a man in a black shirt') map to a massive distribution of possible human faces. Without structural feature embeddings (like SFace or ArcFace vectors), the model samples a random face from noise every time.

Q3: Why did your initial face-overlay approach produce a cutout artifact?

Answer: The initial approach cropped an oval around the user's face and used OpenCV seamlessClone to paste it onto an AI body. This failed because 2D image pasting cannot reconcile 3D lighting differences, skin tone boundaries, or head angles. The true fix was using model-level identity conditioning (InstantID IdentityNet) which synthesizes the entire image from noise conditioned on the face embedding.

Q4: What is the difference between YuNet and SFace in your vision pipeline?

Answer: YuNet is a lightweight face detection ONNX model that outputs bounding boxes and 5 key facial landmarks (eyes, nose, mouth) to calculate pose alignment. SFace is a deep recognition ONNX network that receives the aligned 112x112 face and outputs a 128-dimensional L2-normalized vector capturing the unique facial geometry.

Q5: How does the system handle multiple reference photos of the same person?

Answer: services/face_pipeline.py extracts the 128-d SFace vector for each uploaded photo, checks intra-photo cosine similarity to ensure they are the same person, and computes an averaged master embedding vector that captures the user's identity under multiple lighting angles.

Possible Viva Questions & Answers (Part 2: Engineering)

Anticipated technical questions on implementation, security, and debugging.

VIVA PREPARATION

Q6: How does the system handle Google Gemini API rate limits (HTTP 429)?

Answer: In services/ai_router.py, I implemented dynamic multi-model rotation. If gemini-2.5-flash reaches its free-tier per-model limit, the router catches the exception and immediately rotates to gemini-3.6-flash, gemini-flash-latest, or gemini-3.7-flash, ensuring uninterrupted uptime.

Q7: How did you solve the Windows Python SSL certificate verification failure?

Answer: Python on Windows maintains an isolated certifi bundle that failed on enterprise/proxy networks. I integrated truststore.inject_into_ssl() at application startup, which transparently hooks Python's SSL context into the native Windows OS CryptoAPI certificate store.

Q8: How do you enforce English visual UI while supporting Telugu voice conversations?

Answer: I implemented a 3-stage sanitize_visual_text() filter in telugu_nlp.py. Spoken Telugu is transcribed and understood natively, and JaganGpt synthesizes spoken Telugu audio via gTTS (lang='te'). However, before any text is rendered to the UI, it is sanitized to Latin/English so 0 Telugu Unicode script appears on screen.

Q9: How do you ensure user API keys are never exposed?

Answer: Keys are loaded into Config from server-side .env files. They are never rendered into HTML templates, never sent in JSON payloads to the browser, and never committed to Git. Key verification performs presence checks only ('FOUND' / 'NOT FOUND') without logging the actual characters.

Q10: What was the single biggest technical lesson you learned during this major project?

Answer: The importance of root-cause architecture over superficial patching. When the face cutout bug occurred, attempting to tweak prompt words or adjust mask thresholds failed. The real solution required understanding diffusion conditioning architectures and refactoring the pipeline from crude image compositing to true multimodal latent guidance.

Conclusion, Demonstration & Key Learnings

Final reflections, project competencies gained, and project demonstration.

CONCLUSION

Project Conclusion & Technical Impact

JaganGpt Studio successfully demonstrates that a student developer can engineer a complete, cohesive multimodal AI workspace that matches the core capabilities of commercial foundation platforms while solving one of their greatest limitations: identity preservation.

Key Competencies Gained:

- 1. **Computer Vision:** Deep Feature Embeddings: Implemented ONNX neural inference for facial landmarks (YuNet) and 128-d cosine vectors (SFace).
- 2. **Generative AI:** Foundation LLM Orchestration: Integrated Google Gemini 3.6 Flash and Veo 3.1 with automated quota rotation.
- 3. **Audio Processing:** Audio Engineering: Built real-time speech capture with interruption cancellation and bilingual synthesis.
- 4. **Systems Engineering:** Full-Stack Web Architecture: Built a responsive, secure Flask application with zero-exposure credential handling.

Live Demonstration & Host Access

ACTIVE LOCAL HOST

<http://localhost:5001>

Accessible on any desktop browser. Demonstrates text chat, document analysis, voice interaction, reference image upload, and identity-preserved 4K generation.

MOBILE LAN ACCESS

<http://192.168.8.115:5001>

Accessible from any mobile phone or tablet on the local Wi-Fi network. Demonstrates responsive touch drawer, mobile camera upload, and mobile voice dictation.

Acknowledgments:

I express my sincere gratitude to my Project Guide, Faculty Members, and the Department Review Committee for their guidance and technical feedback throughout the development of JaganGpt Studio.

THANK YOU! • OPEN FOR QUESTIONS & DEMO